# Baby Tetris First Part Report

ERGIN Seçkin Yağmur and GROUSSON Dylan

November 29, 2025

https://github.com/diaarca/baby-tetris/

# Contents

# 1 Question 1: Discounted MDP For Player

The objective of this part is to choose the optimal policy that maximizes the expected cumulative discounted reward over time for the "player".

**States**
The state is represented by the current configuration of the grid, and the next piece to be placed.
$S := (\text{current grid, next piece to place})$

**Actions**
The actions correspond to a position and orientation to place the piece on the grid.
$A := (\text{position, orientation})$

**Transition Function**
The transition function is defined as $P(s'|s,a) := 1/2$
There are two possible pieces (I or L) as the next piece to place, and considering our state representation, and our assumption of uniform distribution between the pieces, each transition has an equal probability of $1/2$.

**Reward Function and Discount Factor**
The reward function is computed based on the number of lines cleared after placing the piece, and a discount factor $\lambda$ is used to weigh future rewards.
$R(s,a,s') := \text{number of lines cleared and } \lambda \in [0,1]$

## 1.1 Finding Optimal Policy

We use the value iteration method to find the optimal policy that maximizes the expected rewards, iteratively updating the value function for each state based on the Bellman equation until it converges to the optimal value (margin $\epsilon$).

$$V(s) = \max_{a \in A} \sum_{s' \in S} P(s'|s, a) \cdot (R(s, a, s') + \lambda V(s'))$$

Value iteration is better suited for maximizing the expected discounted reward since it computes the optimal value then extracts the policy and easier to implement

## 1.2 Implementation

### 1.2.1 Main Architecture

Based on the our design choices from Section 1, we decided to use an object oriented architecture (C++) for our implementation.

For our structure we have declared 7 classes:

1. *Point*: classical point (x,y) class used for grid coordinates

2. *Field*: the class that contain a boolean grid s.t. grid[i][j] = 1 if the corresponding cell is full in the field

3. *Tromino*: the class that correspond to the baby Tetris pieces (I Piece or L Piece)

4. *State*: the class that contain both a field grid and an incoming Tromino

5. *Action*: that represent an action choosable on a given state (combination for placement and rotation of an incoming Tromino)

6. *Game*: the class that contain the configuration (completed lines scores), the current state and score of a baby Tetris game (the configuration is loaded from a file : *config.txt*)

7. *MDP*: the class that all the necessary methods in order to compute the value iteration over a game

Now that we overviewed the global implementation architecture, we will describe our implementation choices.

### 1.2.2 Implementation Choices

- **Piece teleported to their final position**: in order to make the number of intermediate states when choosing an action, we compute on a state all available actions on this state, then according to the policy we make a choice over these actions and we teleport the incoming piece in the current state directly to its final position (with the right rotation)

- **Reward Function**: The reward function as defined in Section 1 take both the current state $s$ and an available action $a_s$. Here since we can calculate the available actions, we can know in which state $s'$ we're going to arrive and we can compute directly in $s'$ the number of completed lines before remove them right to the game. Then we abstract the reward function by counting the number of complete lines in the resulting state

- **Game Ending**: According to the subject, the game should end whenever a piece exceed the height limit of the grid. Again there, we can compute all available actions on a given state, since we cannot place a Tromino outside the grid, such an action will not appear in the set of available actions. Then the game end when there is no available action for the current game state

- **Max Iteration**: Since the value iteration isn't an exact method, if we fix an $\epsilon$ which is too small, we may iterate for a very long moment. For this reason we fix a maximum number of iteration for the expected value improvement, then we always a value iteration which compute in a reasonable time even if the expected value vector is less qualitative than the $\epsilon$ margin we fixed

## 1.3 Results

Here, we will present multiple examples of execution of our implementation. In these executions we will use fixed values for: width = 4, $\epsilon$ = 0.00000001, maxIteration = 100, probaIPiece = 0.5 and maxGameAction = 10000. Note that we fixed our $\epsilon$ that low to be as close to the optimal policy as possible and that our algorithm is in any case bounded in time by the *maxIteration*. In contrary, we will make evolve: the height $\in \{4, 5\}$ and the discount factor $\lambda \in [0, 1]$.

For each execution we will measure:

1. the iterations of the value iteration algorithm

2. the average expected gain of the optimal policy

3. the final score either at the end of the game or if we reach the maximum number of actions

4. the execution time of the said execution

### 1.3.1 Executions

```
./bin/tetris
Loaded config: [1, 3, 6]

All constants:
width = 4, height = 4, probaIPiece
= 0.5, maxGameAction = 10000
epsilon = 1e-08, maxIteration =
100, lambda = 0.1

 Iteration i = 0 and delta = 6
 Iteration i = 1 and delta = 0.6
 Iteration i = 2 and delta = 0.06
 Iteration i = 3 and delta =
0.00525
 Iteration i = 4 and delta =
0.00045
 Iteration i = 5 and delta = 0

 average over final V 0.412926
 ....
 ....
 **..
 **..

 Game Over! Global score: 7849 in
10000 actions
 make run  3.02s user 0.08s system
92% cpu 3.340 total
```

height: 4 and $\lambda$: 0.1

```
./bin/tetris
Loaded config: [1, 3, 6]

All constants:
width = 4, height = 4, probaIPiece
= 0.5, maxGameAction = 10000
epsilon = 1e-08, maxIteration =
100, lambda = 0.4

 Iteration i = 0 and delta = 6
 Iteration i = 1 and delta = 2.4
 Iteration i = 2 and delta = 0.96
 Iteration i = 3 and delta = 0.336
 Iteration i = 4 and delta = 0.1152
 Iteration i = 5 and delta = 0

 average over final V 0.473608
 ***.
 **..
 *.**
 *.**

 Game Over! Global score: 225 in
285 actions
 make run  2.93s user 0.08s system
93% cpu 3.237 total
```

height: 4 and $\lambda$: 0.4

```
 ./bin/tetris                          ./bin/tetris
 Loaded config: [1, 3, 6]              Loaded config: [1, 3, 6]

 All constants:                        All constants:
 width = 4, height = 5, probaIPiece    width = 4, height = 5, probaIPiece
 = 0.5, maxGameAction = 10000          = 0.5, maxGameAction = 10000
 epsilon = 1e-08, maxIteration =       epsilon = 1e-08, maxIteration =
 100, lambda = 0.1                     100, lambda = 0.9

 Iteration i = 0 and delta = 6         Iteration i = 0 and delta = 6
 Iteration i = 1 and delta = 0.6       Iteration i = 1 and delta = 5.4
 Iteration i = 2 and delta = 0.06      Iteration i = 2 and delta = 4.86
 Iteration i = 3 and delta = 0.006     Iteration i = 3 and delta = 4.374
 Iteration i = 4 and delta =           Iteration i = 4 and delta =
 0.00058125                            3.81358
 Iteration i = 5 and delta = 3.1875    Iteration i = 5 and delta =
 e-05                                  1.88219
 Iteration i = 6 and delta = 0         Iteration i = 6 and delta = 0

 average over final V 0.462044         average over final V 0.676855
 ....                                  ....
 ....                                  ....
 ....                                  ....
 ...*                                  ..*.
 .***                                  .***

 Game Over! Global score: 7864 in      Game Over! Global score: 8513 in
 10000 actions                         10000 actions
 make run  58.26s user 0.32s system    make run  58.20s user 0.34s system
 99% cpu 58.837 total                  99% cpu 58.809 total
```

| | |
|:---:|:---:|
| height: 5 and $\lambda$: 0.1 | height: 5 and $\lambda$: 0.9 |

### 1.3.2   Observations

Over our executions, we can remark first that the $\lambda$ have a huge impact on the optimal policy computed. On the small version of the game ($4 \times 4$) it is really difficult to reach the $maxGameAction$ with $\lambda = 0.4$, since the model will prioritize short term decisions in order to increase more quickly its score. In comparison, with $\lambda = 0.1$, the model will not gain that much per action taken, indeed, it will prioritize safer actions that allow itself to gain more score on a long term perspective.

That result was kind of expected since in a $4 \times 4$ grid, there is no much space for risky moves. Then study of the $4 \times 5$ grid case is really interesting because it shows us that with a bigger grid we can easily take much more risks in order to increase our global score. The $height = 5$ and $\lambda = 0.9$ case illustrates perfectly this because it is now stopped by the $maxGameAction$ with an higher score than the $\lambda = 0.1$ case (since the risky actions can be smoothed on a bigger grid, we simply benefit of the higher score).

In both cases, the total computational time is pretty acceptable (max 1 minute) even if it can be largely optimized (see Section 1.4). About the exact same improvement, the average expected gain $V$ cannot be really discussed here since there is a huge part of states that aren't reachable in practice.

Finally, we can remark that the value iteration algorithm choice described in Section 1 was pretty relevant because we converge very quickly to the optimal policy (at the last iteration, $\Delta$ is always equal to 0).

### 1.3.3  Usage

Now, if you want to try our code by yourself, you can modify the different constants by modifying the corresponding **#define** statements in **src/Tetris.cpp**, **hdr/MDP.h** and **hdr/State.h**.
In order to recompile and run the whole project, you simply need to execute the command:

```
make run
```

Contrarily to the textual output shown in Section 1.3.1, in the version that you will receive, you will have an exhaustive description of the played game (evolution of the state at each action) as follows:

```
....                  .*..                  .*..
..*.                  ***.                  ***.
***.  --- LPiece -->  ***.  -- Completion -> ***.
***.                  ***.                  ***.
Current  score:  734

.*..                  .***                  ....
***.                  ****                  .***
***.  --- LPiece -->  ***.  -- Completion -> ***.
***.                  ***.                  ***.
Current  score:  735

....                  ***.                  ***.
.***                  .***                  .***
***.  --- IPiece -->  ***.  -- Completion -> ***.
***.                  ***.                  ***.
Current  score:  735
```

## 1.4   Incoming Improvement

For now we compute naively the value iteration algorithm on all state combinations, which represents

$$= 2^{\text{gridHeight}*\text{gridWidth}} * \text{nbOfTrominoTypes}$$
$$= 2^{4*4} * 2$$
$$= 2^{17}$$
$$= 131072$$

states in the smallest $(4 \times 4)$ version of the game. In the near future we want to considerably reduce this number by computing the value iteration only an states that are reachable from a given starting state $s_0$. For this purpose, we will make use of a Breadth First Search algorithm which will explore iteratively all reachable states from the current one.